

```
In [31]: import numpy as np
import os
from sklearn.metrics import confusion_matrix
import seaborn as sn; sn.set(font_scale=1.4)
from sklearn.utils import shuffle
import matplotlib.pyplot as plt
import cv2
import tensorflow as tf
from tqdm import tqdm
```

WARNING:tensorflow:From C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
In [5]: class_names = ['mountain', 'street', 'glacier', 'buildings', 'sea', 'forest']
class_names_label = {class_name:i for i, class_name in enumerate(class_names)}

nb_classes = len(class_names)

print("class_names_label", class_names_label)
print("nb_classes", nb_classes)

IMAGE_SIZE = (150, 150)
```

```
class_names_label {'mountain': 0, 'street': 1, 'glacier': 2, 'buildings': 3, 'sea':
4, 'forest': 5}
nb_classes 6
```

```
In [32]: def load_data():
        """
        Load the data:
        - 14,034 images to train the network.
        - 3,000 images to evaluate how accurately the network learned to classify
        """

        datasets = [r'C:\Users\PC\OneDrive\Desktop\seg_train', r'C:\Users\PC\OneDrive\Desktop\seg_test']
        output = []

        # Iterate through training and test sets
        for dataset in datasets:

            images = []
            labels = []

            print("Loading {}".format(dataset))

            # Iterate through each folder corresponding to a category
            for folder in os.listdir(dataset):
                label = class_names_label[folder]

                # Iterate through each image in our folder
                for file in tqdm(os.listdir(os.path.join(dataset, folder))):

                    # Get the path name of the image
                    img_path = os.path.join(os.path.join(dataset, folder), file)
```



```
Number of training examples: 14034
Number of testing examples: 3000
Each image is of size: (150, 150)
(14034,)
(3000,)
```

```
In [36]: train_images = train_images / 255.0
        test_images = test_images / 255.0
```

```
In [37]: def display_random_image(class_names, images, labels):
        """
            Display a random image from the images array and its correspond label from the
        """

        index = np.random.randint(images.shape[0])
        plt.figure()
        plt.imshow(images[index])
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)
        plt.title('Image #{0} : '.format(index) + class_names[labels[index]])
        plt.show()
```

```
In [38]: display_random_image(class_names, train_images, train_labels)
```

Image #805 : buildings



```
In [40]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

        # Create an instance of the ImageDataGenerator
        datagen = ImageDataGenerator(
            rotation_range=40,
            width_shift_range=0.2,
            height_shift_range=0.2,
```

```

shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,
fill_mode='nearest'
)

# Visualize a few augmented images
example_image = train_images[0]
example_image = example_image.reshape((1,) + example_image.shape) # Reshape to (1, height, width, channels)

fig, ax = plt.subplots(1, 5, figsize=(15, 3))

for i, augmented_image in enumerate(datagen.flow(example_image, batch_size=1)):
    ax[i].imshow(augmented_image[0])
    ax[i].axis('off')

    if i == 4:
        break

plt.show()

```



```

In [41]: # Create an instance of the ImageDataGenerator with augmentation parameters for training
train_datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

# Create an instance of the ImageDataGenerator without augmentation for validation
val_datagen = ImageDataGenerator()

```

```

In [42]: # Define callbacks
early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
model_checkpoint = tf.keras.callbacks.ModelCheckpoint('best_model.h5', save_best_only=True)

```

```

In [43]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Conv2D(128, (3, 3), activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(1000, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

```

```
tf.keras.layers.BatchNormalization(),
tf.keras.layers.MaxPooling2D(2, 2),

tf.keras.layers.Conv2D(256, (3, 3), activation='relu'),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.MaxPooling2D(2, 2),

tf.keras.layers.Conv2D(512, (3, 3), activation='relu'),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.MaxPooling2D(2, 2),

tf.keras.layers.Flatten(),

tf.keras.layers.Dense(128, activation='relu'),
tf.keras.layers.Dropout(0.5),

tf.keras.layers.Dense(6, activation='softmax')
])
model.summary()
```

WARNING:tensorflow:From C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 32)	896
batch_normalization (Batch Normalization)	(None, 148, 148, 32)	128
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 72, 72, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 34, 34, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 256)	295168
batch_normalization_3 (Batch Normalization)	(None, 15, 15, 256)	1024
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 256)	0
conv2d_4 (Conv2D)	(None, 5, 5, 512)	1180160
batch_normalization_4 (Batch Normalization)	(None, 5, 5, 512)	2048
max_pooling2d_4 (MaxPooling2D)	(None, 2, 2, 512)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262272
dropout (Dropout)	(None, 128)	0

dense_1 (Dense) (None, 6) 774

```
=====
Total params: 1835590 (7.00 MB)
Trainable params: 1833606 (6.99 MB)
Non-trainable params: 1984 (7.75 KB)
```

```
In [44]: # Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model with data augmentation for training data and without augmentation for
history = model.fit(
    train_datagen.flow(train_images, train_labels, batch_size=32),
    steps_per_epoch=len(train_images) // 32,
    epochs=20,
    validation_data=val_datagen.flow(test_images, test_labels, batch_size=32),
    validation_steps=len(test_images) // 32,
    callbacks=[early_stopping, model_checkpoint]
)
```

WARNING:tensorflow:From C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Epoch 1/20

WARNING:tensorflow:From C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

438/438 [=====] - 426s 958ms/step - loss: 1.4518 - accuracy: 0.4814 - val_loss: 1.9771 - val_accuracy: 0.3353

C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\engine\training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
 saving_api.save_model(

Epoch 2/20

438/438 [=====] - 415s 947ms/step - loss: 1.1190 - accuracy: 0.5781 - val_loss: 1.1724 - val_accuracy: 0.5736

Epoch 3/20

438/438 [=====] - 407s 928ms/step - loss: 1.0214 - accuracy: 0.6193 - val_loss: 0.9157 - val_accuracy: 0.6384

Epoch 4/20

438/438 [=====] - 413s 942ms/step - loss: 0.9486 - accuracy: 0.6470 - val_loss: 1.1902 - val_accuracy: 0.5501

Epoch 5/20

438/438 [=====] - 381s 869ms/step - loss: 0.8862 - accuracy: 0.6770 - val_loss: 1.0246 - val_accuracy: 0.6442

Epoch 6/20

438/438 [=====] - 384s 876ms/step - loss: 0.8543 - accuracy: 0.6913 - val_loss: 2.2552 - val_accuracy: 0.4550


```
In [45]: def plot_accuracy_loss(history):
        """
        Plot the accuracy and the loss during the training of the nn.
        """
        fig = plt.figure(figsize=(10,5))

        # Plot accuracy
        plt.subplot(221)
        plt.plot(history.history['accuracy'], 'bo--', label = "acc")
        plt.plot(history.history['val_accuracy'], 'ro--', label = "val_acc")
        plt.title("train_acc vs val_acc")
        plt.ylabel("accuracy")
        plt.xlabel("epochs")
        plt.legend()

        # Plot Loss function
        plt.subplot(222)
        plt.plot(history.history['loss'], 'bo--', label = "loss")
        plt.plot(history.history['val_loss'], 'ro--', label = "val_loss")
        plt.title("train_loss vs val_loss")
        plt.ylabel("loss")
        plt.xlabel("epochs")

        plt.legend()
        plt.show()
```

```
In [46]: plot_accuracy_loss(history)
```




```
In [48]: # Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_images, test_labels, verbose=2)
print("\nTest Accuracy:", test_accuracy)

# Make predictions on the test set
predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)

# Display random images with actual and predicted labels
def display_random_images(images, actual_labels, predicted_labels, class_names, num_images):
    indices = np.random.choice(len(images), num_images, replace=False)

    plt.figure(figsize=(15, 5))
    for i, index in enumerate(indices, 1):
        plt.subplot(1, num_images, i)
        plt.imshow(images[index])
        plt.title(f"Actual: {class_names[actual_labels[index]]}\nPredicted: {class_names[predicted_labels[index]]}")
        plt.axis('off')

    plt.show()

# Display random images
display_random_images(test_images, test_labels, predicted_labels, class_names)
```

94/94 - 19s - loss: 0.9134 - accuracy: 0.6390 - 19s/epoch - 206ms/step

Test Accuracy: 0.6389999985694885

94/94 [=====] - 20s 212ms/step


Actual: buildings
Predicted: buildings



Actual: street
Predicted: street




Actual: sea
Predicted: sea



Actual: glacier
Predicted: glacier




Actual: sea
Predicted: sea



```
In [49]: from sklearn.metrics import confusion_matrix
import seaborn as sns

# Generate the confusion matrix
conf_matrix = confusion_matrix(test_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.show()
```



```
In [50]: # Load the pre-trained InceptionV3 model without the top layers
base_model = tf.keras.applications.InceptionV3(weights='imagenet', include_top=False,

# Choose the number of layers to freeze
num_layers_to_freeze = 20

# Freeze only the first num_layers_to_freeze layers
for layer in base_model.layers[:num_layers_to_freeze]:
    layer.trainable = False
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
 87910968/87910968 [=====] - 37s 0us/step

```
In [51]: # Import models module from Keras
from tensorflow.keras import models
# Import necessary layers module from Keras
from tensorflow.keras import layers

# Create a new model by adding custom layers on top of VGG16
model = models.Sequential([
    base_model,
    layers.Flatten(),
    layers.Dense(256, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(128, activation='relu'),
```

```

layers.Dropout(0.5),
layers.Dense(nb_classes, activation='softmax')
])

```

```

In [52]: # Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

```

In [53]: # Create an instance of the ImageDataGenerator with augmentation parameters for training
train_datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

# Create an instance of the ImageDataGenerator without augmentation for validation
val_datagen = ImageDataGenerator()

```

```

In [54]: # Define callbacks
early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3, rest_model_checkpoint = tf.keras.callbacks.ModelCheckpoint('best_model_vgg16.h5', save_best

```

```

In [55]: # Train the model with data augmentation for training data and without augmentation for validation
history = model.fit(
    train_datagen.flow(train_images, train_labels, batch_size=32),
    steps_per_epoch=len(train_images) // 32,
    epochs=20,
    validation_data=val_datagen.flow(test_images, test_labels, batch_size=32),
    validation_steps=len(test_images) // 32,
    callbacks=[early_stopping, model_checkpoint]
)

```

Epoch 1/20

438/438 [=====] - ETA: 0s - loss: 1.4865 - accuracy: 0.3809

C:\Users\PC\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\engine\training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.

saving_api.save_model(

```

438/438 [=====] - 805s 2s/step - loss: 1.4865 - accuracy: 0.
3809 - val_loss: 1.0766 - val_accuracy: 0.4929
Epoch 2/20
438/438 [=====] - 856s 2s/step - loss: 1.0276 - accuracy: 0.
5681 - val_loss: 0.8652 - val_accuracy: 0.6418
Epoch 3/20
438/438 [=====] - 743s 2s/step - loss: 0.9421 - accuracy: 0.
6298 - val_loss: 0.9575 - val_accuracy: 0.6183
Epoch 4/20
438/438 [=====] - 759s 2s/step - loss: 0.8713 - accuracy: 0.
6603 - val_loss: 1.4192 - val_accuracy: 0.4677
Epoch 5/20
438/438 [=====] - 822s 2s/step - loss: 0.9250 - accuracy: 0.
6607 - val_loss: 0.7869 - val_accuracy: 0.6962
Epoch 6/20
438/438 [=====] - 820s 2s/step - loss: 0.7979 - accuracy: 0.
7418 - val_loss: 0.7433 - val_accuracy: 0.7806
Epoch 7/20
438/438 [=====] - 817s 2s/step - loss: 0.7492 - accuracy: 0.
7463 - val_loss: 0.6401 - val_accuracy: 0.7409
Epoch 8/20
438/438 [=====] - 941s 2s/step - loss: 0.7010 - accuracy: 0.
7608 - val_loss: 0.5687 - val_accuracy: 0.8017
Epoch 9/20
438/438 [=====] - 763s 2s/step - loss: 0.7042 - accuracy: 0.
7594 - val_loss: 1.1795 - val_accuracy: 0.6109
Epoch 10/20
438/438 [=====] - 765s 2s/step - loss: 0.6153 - accuracy: 0.
7930 - val_loss: 0.4036 - val_accuracy: 0.8606
Epoch 11/20
438/438 [=====] - 1200s 3s/step - loss: 0.5841 - accuracy:
0.8034 - val_loss: 0.6281 - val_accuracy: 0.7860
Epoch 12/20
438/438 [=====] - 763s 2s/step - loss: 0.6277 - accuracy: 0.
8074 - val_loss: 0.5683 - val_accuracy: 0.7876
Epoch 13/20
438/438 [=====] - 769s 2s/step - loss: 0.5920 - accuracy: 0.
8049 - val_loss: 0.4336 - val_accuracy: 0.8501

```

```

In [56]: # Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_images, test_labels, verbose=2)
print("\nTest Accuracy:", test_accuracy)

# Make predictions on the test set
predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)

# Display random images with actual and predicted labels
def display_random_images(images, actual_labels, predicted_labels, class_names, num_in
    indices = np.random.choice(len(images), num_images, replace=False)

    plt.figure(figsize=(15, 5))
    for i, index in enumerate(indices, 1):
        plt.subplot(1, num_images, i)
        plt.imshow(images[index])
        plt.title(f"Actual: {class_names[actual_labels[index]]}\nPredicted: {class_nam
        plt.axis('off')

```

```
plt.show()
```

```
# Display random images
```

```
display_random_images(test_images, test_labels, predicted_labels, class_names)
```

94/94 - 48s - loss: 0.4013 - accuracy: 0.8617 - 48s/epoch - 514ms/step

Test Accuracy: 0.8616666793823242


94/94 [=====] - 39s 395ms/step



```
In [57]: from sklearn.metrics import confusion_matrix
import seaborn as sns

# Generate the confusion matrix
conf_matrix = confusion_matrix(test_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, y
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.show()
```



```
In [ ]:
```